
Communication-Efficient & Class-Balancing Federated Learning with Adaptive Dropout & Model Quantization

Anonymous Author(s)

Affiliation

Address

email

Abstract

Federated learning (FL) is used to train deep learning models over heterogeneous and decentralized datasets. Communication between client and server can be a major bottleneck for Federated learning, specially in case of large models. Moreover there is often data imbalance issue in real world FL problems.

To address the above mentioned issues we propose a novel balanced communication efficient federated learning strategy. Our strategy is based on two key strategies : 1. A novel Adaptive federated dropout on client models to address both communication bottleneck and data imbalance. 2. Quantization method to improve communication efficiency. We empirically show that the combination of these two strategies improve communication strategy while improving the accuracy of the model on several data set and experimental settings. The code can be found at anonymous.

1 Introduction

1.1 Motivation for Federated Learning

The digital era is witnessing an unprecedented increase in data generation, much of which is highly sensitive and personal. This surge necessitates the advancement of deep learning models that can be trained over decentralized and heterogeneous datasets, thus ushering in the era of federated learning (FL). FL represents a paradigm shift from traditional centralized machine learning approaches by enabling the training of models directly on devices where the data is generated (30). This method ensures data privacy since the raw data never leaves its original device, thereby addressing growing concerns over data security and privacy in the age of information (24; 47).

The cornerstone of FL lies in its ability to maintain the confidentiality and integrity of data. In traditional machine learning approaches, data aggregation poses significant risks of privacy breaches. However, FL aims to circumvent this issue by decentralizing the data processing, thereby safeguarding sensitive information. This approach can not only enhance user trust but also complies with stringent data protection regulations such as the General Data Protection Regulation (GDPR) (41) in Europe.

While FL presents an innovative solution to privacy concerns, it introduces unique challenges stemming from the decentralized nature of data. These challenges include managing data heterogeneity across different devices and ensuring model performance despite the varying quality and quantity of local datasets. The decentralized framework of federated learning necessitates novel approaches to model training and aggregation to tackle these obstacles efficiently while maintaining the aforementioned privacy tenets (19).

33 1.2 Challenges in Federated Learning

34 1.2.1 Communication Bottleneck

35 One of the paramount challenges in FL is the communication bottleneck that arises, especially with
36 large deep learning models. The necessity to frequently exchange model updates between the server
37 and numerous clients can lead to significant communication overhead (28). This issue is exacerbated
38 in scenarios with limited bandwidth or when the model size becomes substantially large, thereby
39 straining network resources and impacting the efficiency of the learning process.

40 1.2.2 Data Imbalance

41 Another significant challenge is the inherent data imbalance across clients in FL settings. Given
42 that data is generated in a non-IID (independent and identically distributed) manner across different
43 devices, some clients may have more data or more diverse data than others. This imbalance can lead
44 to skewed model updates, where the model disproportionately learns from clients with more or more
45 varied data, potentially compromising the overall model performance and fairness (46; 13).

46 1.2.3 Proposed Solution

47 Neural Network quantization techniques can play a critical role in mitigating communication bot-
48 tlenecks. While regularization methods like dropout can improve model performance as well as
49 improve communication efficiency (39). Dropout works by randomly "dropping out" (setting to
50 zero) a certain percentage of neural network units during the training. When combined with an
51 appropriate quantization method dropout can significantly reduce the amount of data transmission
52 between client and server. Such reduction is vital for the efficient transmission of model updates
53 over networks with limited bandwidth (1). Furthermore, incorporating dropout can also serve as an
54 effective regularization technique. It improves model robustness against over-fitting and enhances its
55 generalization capabilities across decentralized datasets (39?).

56 Furthermore, the application of dropout and quantization techniques can inadvertently bolster the
57 security of federated learning systems. By reducing the model's complexity, the potential attack
58 surface for adversarial attacks is minimized, thereby safeguarding both the integrity of the model and
59 the privacy of user data (40). This dual benefit underscores the importance of model compression
60 not only for improving FL efficiency and scalability but also for reinforcing the security posture of
61 federated learning ecosystems against evolving cyber-security threats.

62 In this paper, we propose two novel strategies to mitigate the communication footprint in federated
63 learning and provide empirical evidence demonstrating that our model is robust to data imbalance.
64 The specific contributions of this study are as follows:

65 **Adaptive Federated Dropout** innovates by integrating the conventional Dropout (39) technique,
66 previously proposed federated dropout (43; 2) with and a voting based adaptive dropout method
67 inspired by CLIMB Algorithm (35). Our method involves training client models that start with a
68 fixed dropout rate, but adjust this rate dynamically based on model performance feedback from the
69 server. The clients vote to either increase or decrease the dropout and based on that server makes a
70 decision. If a client model achieves higher accuracy than what the server model returns, then it votes
71 to increase the dropout rate; conversely, it votes to decrease if the server's model performs better.

72 **Model Quantization** forms our second approach to reducing the communication footprint in federated
73 learning environments significantly. By implementing various quantization levels introduced before
74 (12), from float16 to int8, this method efficiently lowers bandwidth requirements while negligibly
75 affecting model performance. Our experiments across multiple benchmark datasets showcase the
76 robustness of model quantization, particularly when combined with adaptive dropout techniques,
77 even on datasets with class imbalances like imbalanced-cifar.

78 **Empirical Validation** through experiments on several benchmark datasets confirms the effectiveness
79 of our proposed strategies. The combined use of adaptive dropout and model quantization not only
80 conserves bandwidth but also proves robust against class imbalances, thereby providing a compelling
81 case for their application in practical federated learning scenarios. We also provide our source-code
82 with examples to make it easy to integrate our method with existing federated learning frameworks
83 and different model architectures.

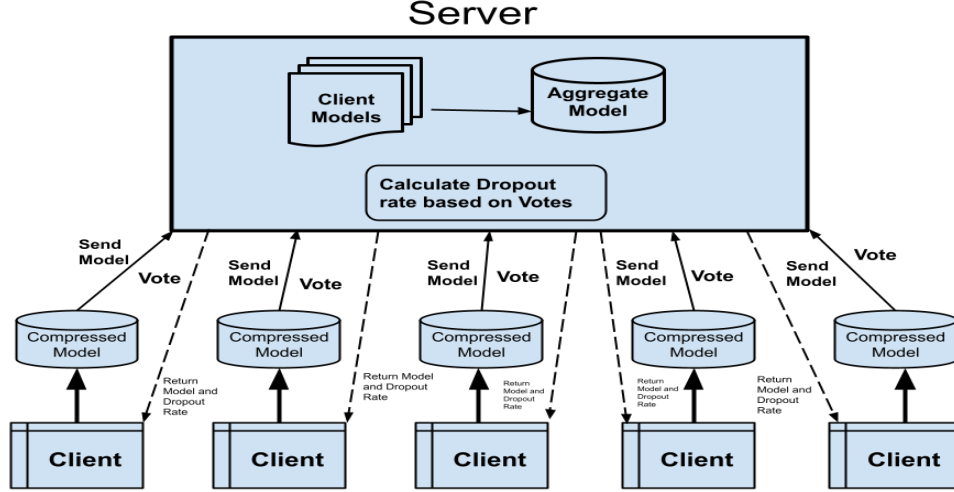


Figure 1: A block diagram of Adaptive Federated learning with Quantization.

2 Related Work

2.1 Foundations of Federated Learning

Federated learning (FL) has emerged as a paradigm that leverages vast, dispersed datasets while ensuring privacy protection. Originating from the work of McMahan et al.(25) in 2016, FL enables model training across multiple devices without centralizing data, thus addressing the privacy and security limitations of traditional cloud-based methods. By decentralizing the process, FL enhances privacy and reduces the bandwidth needed for data transfer, a significant advantage in the era of rapid digital expansion (8).

At the core of FL are principles of collaborative model training. Data remains on local devices like smartphones and IoT devices, and only model updates are shared with a central server. This server aggregates the updates to improve the global model, which is then distributed back to the devices for further training (20). Through iterative cycles, the model enhances using diverse data sources while maintaining data privacy, aligning with strict regulations such as GDPR (41).

A key development in FL was the creation of the FederatedAveraging algorithm (FedAvg) by McMahan et al (26), which reduces communication overhead and speeds up model convergence by averaging local updates, optimizing both efficiency and effectiveness.

Subsequent FL research has tackled challenges like non-IID data, system heterogeneity, and defense against adversarial attacks. Notable contributions include Konečný et al(15) methods for efficient communication.

2.2 Communication Efficiency and Quantization

FL signifies a substantial shift in machine learning by decentralizing data processing, thus enhancing data privacy and security. This transformation introduces a critical challenge: ensuring communication efficiency. Aggregating updates from numerous devices to a central server can be bandwidth-intensive and slow, particularly as machine learning models grow in complexity (?).

To achieve communication efficiency, FL strategies focus on reducing bandwidth and latency during large-scale model update transmissions (18). Essential methods include quantization, pruning, and compression techniques, all aimed at decreasing the data volume communicated. Quantization, which

reduces the precision of model parameters, significantly lessens data transmission volume while only marginally affecting model accuracy (12).

Quantization in federated learning has evolved from basic probabilistic approaches for model weight reduction (15) to more sophisticated methods. Initial research such as that by Benoit et al. (12) in 2018 highlighted the limitations of early over-parameterized proofs of concept. Recent methodologies, like QuPeD (32), innovate by integrating distillation with vector quantization, enhancing precision over scalar alternatives. Further exploration into vector quantization has been pursued through various frameworks including UVeQFed (36), HeteroSAG (5), JoPEQ (17), and the latest, FedVQCS (31). Additionally, approaches such as T-FedAvg (44) utilize ternary quantization to simplify weights to -1, 0, 1. Innovations continue with non-uniform (3) and stochastic quantization, which employs random rounding to mitigate the constraints of traditional deterministic methods (21).

2.3 Addressing Data Imbalance

Data imbalance refers to scenarios where data across various nodes are unevenly distributed, either in quantity or in quality. This uneven distribution can significantly skew the learning model, as nodes with fewer data samples—often containing minority representations—might be underrepresented in the training of the global model. Consequently, this leads to biased or inadequate generalizations when the model is subsequently applied to unseen data. Such imbalances pose a critical challenge in developing robust machine learning models that perform well across diverse datasets.

Data Distribution Management encompasses strategies that adjust the distribution of data across nodes to enhance model equity. Techniques include Re-balancing Techniques and Transfer Learning (33; 23). Algorithmic modifications aim to refine learning processes to handle imbalanced data better by modifying optimization objectives to emphasize minority classes, integrating fairness to reduce biases across nodes (4), and implementing regularization methods such as dropout to mitigate over-fitting and manage disparities in data (?). Additional methodologies, such as anomaly detection, re-calibrate focus towards underrepresented data (37), while inherently adaptive strategies like meta-learning(10; 11) and multi-task learning(27) are tailored to optimize performance based on the unique distributions at each node.

2.3.1 Client Dropout

Client dropout, also known as device dropout, poses a challenge to the validity of models in federated learning. This phenomenon primarily affects participant devices, such as IoT devices or smartphones, each treated as an individual node. Dropout at this level can introduce substantial biases, as it may exclude crucial data subsets from the training process, thereby skewing results and compromising the model’s reliability. Extensive research efforts focus on addressing this form of dropout to ensure the integrity of training outcomes (24; 22; 34).

2.4 Dropout Techniques

Dropout in the traditional definition has likewise developed many avenues of research. Some models fall under the type of global dropout, which is where all clients have the same technique applied to them. The other side is called local dropout, or sometimes sub-model learning (14). Local dropout is where partitions of the network are given differential parameters for dropout. This differentiation is an operational one and often times whether in the same study or different ones, the same techniques are often applied in both the global and local form.

Some techniques include using randomness as part or most of the process(43; 2), FedDrop is notable example. Other techniques use Bayesian Inference-based Dropout (7; 45), FjORD (9) is a seminal example, dropping partial back rows in weight matrices, which places pike-and-slab distributions over weights. Other techniques us design score maps like AFD (28), an early adaptive technique. Still others use differential-parameter dropout methods like FedDD (6; 28), these are purpose built with local dropout in mind. Some models build on these by parameterizing based on the the compute power of the device (43; 42), though the nature of this may infringe on privacy of the client, a tenet of FL. A technique building off of those is called FedSpu (29) which deferentially applies parameters under similar conditions of compute ability, but instead of pruning nodes from layers, it freeze gradients during back-propagation.

3 Method

In this section, we present strategies for reducing Federated Learning’s communication costs, namely quantization techniques and Adaptive Federated Dropout. These two strategies are described separately, but are fully compatible.

3.1 End-of-Round Quantization

Our approach to model compression involves End-of-Round Quantization, which optimizes the precision of model parameters at the conclusion of training rounds. This method not only reduces the model’s memory footprint but also conserves bandwidth during data transmission between clients and the server.

Algorithm 1 Adaptive Dropout and Model Quantization in Federated Learning

```

1: Input: Clients  $C = \{C_1, C_2, \dots, C_n\}$ , Server  $S$ , Data  $D_i$  in each client, Client epochs per round  $e$ , Number of rounds  $R$ , Initial dropout rate  $dr$ ,  $v_{acc}$  for the server model validation accuracy,  $l_{acc_i}$  for each client last validation accuracy.
2: Output: A trained model on  $S$ 
3: procedure FEDLR( $C, S, R$ )
4:    $model.setDropoutRate(dr)$ 
5:   for  $r = 1$  to  $R$  do
6:     for each client  $C_i$  in  $C$  do
7:        $W_i \leftarrow \text{Client\_train}(C_i, S.getModel(), dr)$ 
8:        $l_{acc_i} = \text{evaluate}(C_i, D_i)$ 
9:        $v_i \leftarrow \text{vote}(v_{acc}, l_{acc_i})$ 
10:       $QW_i \leftarrow \text{quantizeModel}(W_i, Q)$ 
11:      Send  $QW_i, v_i$  to server
12:    end for
13:     $S.ServerUpdate(\text{aggregate}(QW_1, QW_2, \dots, QW_n))$ 
14:     $S.adjustDropoutRate(v_1, v_2, \dots)$ 
15:    send  $dr$  to clients.
16:  end for
17: end procedure
18: procedure VOTE( $v_{acc}, l_{acc_i}$ )
19:   if  $v_{acc} > l_{acc_i}$  then
20:      $v_i = 1$ 
21:   else
22:      $v_i = -1$ 
23:   end if
24:   return  $v_i$ 
25: end procedure
26: function ADJUSTDROPOUTRATE( $v_1, v_2, \dots$ )
27:    $r = (\sum v_i)/n$ 
28:   if  $r > 0$  then
29:      $dr \leftarrow \min((1 - \alpha), dr(1 + \beta))$ 
30:   else
31:      $dr \leftarrow \max(\alpha, dr(1 - \beta))$ 
32:   end if
33:   return  $dr$ 
34: end function

```

3.2 Adaptive Dropout

To further reduce communication cost we introduce adaptive federated dropout. Our algorithm is inspired by previously proposed works on federated dropout(43; 28; 2) and Shen et al. (35) work on handling class imbalance. Dropout is a popular method of regularization in deep learning where a set of neurons are randomly dropped. In traditional dropout, hidden units are multiplied by a random binary masks with a specified size to drop some neurons during the training. In federated dropout

177 instead of training with one global network, clients train a sub-network based on a dropout rate and
 178 sends it to the server. The server then maps back to the global model, aggregates these models,
 179 performs it's update step and then sends a model for clients to continue training.

180 We propose a dropout strategy where the clients start with a specified dropout rate and then in each
 181 round votes to either increase or decrease the dropout rate. The server aggregates the votes and based
 182 on that either increases or decrease the dropout rate and then communicates the new dropout rate
 183 to all the clients. Clients voting is dependent on it's data distribution. If the accuracy of a client
 184 model surpasses that of the server model returned, then it opts to vote for increasing the dropout
 185 rate; otherwise, it chooses to lower the dropout rate if the server model outperforms. This acts as a
 186 regularization based on the local data distribution while still being applicable to the general FedAvg
 187 algorithm.

188 To be more precise, we start training with n clients each with Data set D_i to be trained for epochs e , a
 189 server S and an initial dropout rate of dr . Client devices save the initial validation accuracy as l_{acc} ,
 190 train a model for e epochs and send it to the server. After each round the server S would aggregate
 191 client models using FedAvg and return the average model to the client. Now the client would have a
 192 new model with a validation accuracy v_{acc} . Each client sends a vote v_i based on the following formula
 193 : If $v_{acc} > l_{acc}$, then vote $\{v_i = 1\}$, If $v_{acc} < l_{acc}$, then vote $\{v_i = -1\}$. The server then aggregates
 194 all the vote values in result, $r = (\sum v_i)/n$. If r is more than 0 then the dropout rate is adjusted as
 195 $dr \leftarrow \min((1 - \alpha), dr(1 + \beta))$ and if it's less than 0 then it's adjusted as $dr \leftarrow \max(\alpha, dr(1 - \beta))$.
 196 This introduces two more hyper-parameters namely α to limit the amount of dropout and β to specify
 197 the amount of adaptation to the dropout rate. We recommend a value of .05 for β and .1 for α . Based
 198 on those values the dropout rate would keep increasing or decreasing by .05 until they are equal to
 199 .1 or .9. The detailed steps of our algorithm are given below in Algorithm 34 This method ensures
 200 that the an adaptive dropout rate is applied in each round based on the data distribution in clients
 201 and the server is still able to aggregate the models with reduces size. This improves communication
 202 efficiency dues to clients sending updates based on smaller sub-networks. The adaptation of the
 203 dropout rate is based on the individual clients data distribution and that makes it more robust to data
 204 imbalance. This technique has a few additional advantages. It makes client training more efficient
 205 due to clients training on smaller sub-networks. Other than that every dropout is different due to
 206 randomness and that introduces a degree of penalization for each client.

207 4 Results

208 In this section, we present our experimental setup, data set, different hyper-parameter used along with
 209 results for using adaptive dropout and quantization together. We limit our experiments to established
 210 benchmark of Federated Averaging (FedAvg). We apply our proposed Adaptive Federated Dropout
 211 along with 'float16' quantization and compare the results with FedAvg base.

212 4.1 Data sets

213 **CIFAR-10 & CIFAR-100:** CIFAR-10 (16) consists of 60,000 color images of 32x32 pixels, or
 214 up to 64x64 pixels, distributed equitably across ten categories, including animate and inanimate
 215 subjects. CIFAR-100 (38), which is even more challenging, contains the same number of images
 216 scaled similarly but divided into 100 categories, making it a more detailed classification task. These
 217 data-sets serve as primary benchmarks. CIFAR-10 encompasses 60,000 32x32 pixel color images.
 218 CIFAR-100 comprises of an identical count of images, but spans over 100 classes.

219 **Model:** We use a CNN model consisting of one (3×3) convolution layer, followed by a one max-pool
 220 and several fully connected layers. We use softmax as an activation in the final layer. For Cifar data
 221 sets we use an input size of $(32 \times 32 \times 3)$. This model has a total of 923k trainable parameters. We
 222 use validation accuracy as evaluation method. We use TensorFlow for implementation of the models.

223 **Hyper-parameters:** We have used Adam optimizer with sparse categorical cross-entropy as loss
 224 function. We use the known baseline settings of FedAvg without any changes. Our federated learning
 225 experiments are performed with 10 clinets running simultaneously using the Flower framework.

226 **Quantization:** Our quantization is inspired by previous work of Jacob et al (12). We use 'float16'
 227 quantization which is able to achieve 2x compression with minimal accuracy loss. We apply
 228 quantization at the end of each client round.

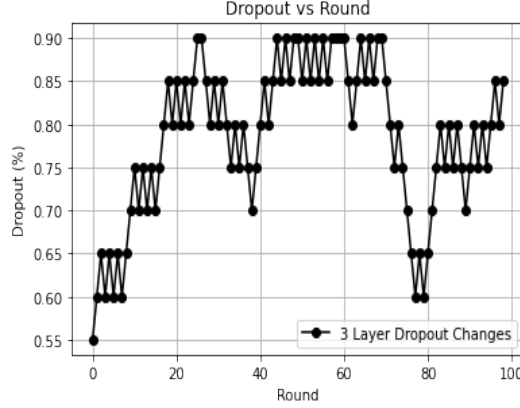


Figure 2: Change in dropout rate throughout the training process.

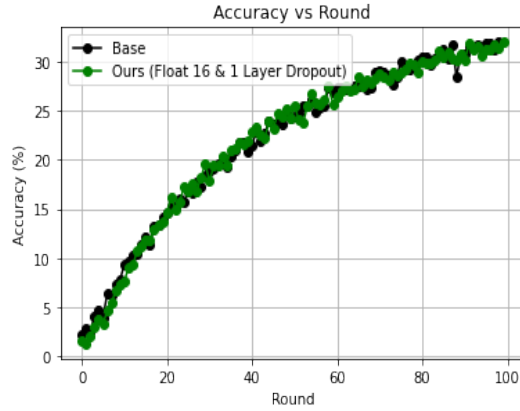


Figure 3: Accuracy of our method compared with base FedAvg on Cifar100. The plot shows that our validation accuracy is better the base case.

Adaptive federated dropout: For base case our model doesn't have any dropout. We compare this with adaptive dropout applied after 1 fully connected layer and applied after 3 fully connected layers. We only experiment with applying dropout on fully-connected layers, since CNN layers have very few parameters. Our adaptive federated dropout has two more hyper-parameters α and β . We start with a dropout rate of .5 and use .05 as value of β , .1 as value of α . That means we increase or decrease the dropout rate by .05, unless it's equal to .9 or .1. We plot the changes in dropout rate throughout the training process in Figure 2.

We plot the results of Cifar 10 data set with 'float16' quantization (2x compression) and different amount of dropout in Figure 4. A dropout in one layer combined with quantization leads to approximately 3x reduction in communicated parameter size and a 3 layers of dropout with quantization leads to approximately 4x reduction in parameter size. The results for Cifar100, presented in Figure 3 show that our method achieves similar results accross different data sets.

Robustness to class Imbalance: To test the robustness to imbalanced data we use imbalanced Cifar10 with the number of samples for class 0 and 1 reduced by half. We plot the results of imbalanced cifar10 in Figure 5. The results show that our method provides advantage over FedAvg in case of imbalanced data set. The dropout rate increase in clients with more data from majority classes and works as a regularization.

Limitations Our study assumes that all node devices involved in the federated learning network are capable of meeting the computational requirements necessary for efficient algorithm performance. This assumption might not accurately reflect the varied capabilities of devices in practical, real-world applications and could skew the model's performance metrics. To counteract the risk of over-fitting,




Figure 4: Accuracy of our method compared with base FedAvg on Cifar10. The plot shows that our validation accuracy is very similar to the base case even when we remove significant amount of parameters with dropout and quantization.




Figure 5: Accuracy of our method compared with base FedAvg on Imbalanced Cifar10. The plot shows that our provides advantage over FedAvg in case of class imbalance while significantly reducing communication needs.

we have employed multiple datasets, namely datasets Cifar10 and Cifar100. Utilizing multiple datasets helps to verify that our model’s efficacy is generalizable and not overly tailored to specific data characteristics. Using quantization would lead to some loss of precision and that might introduce domain adaptation problem in some cases.

5 Conclusion

Federated learning is dependent on client device performance and communication bandwidth. Both of these can vary a lot based on different demographic groups and thereby exclude people from training models. Research on dropout and quantization methods can address this problem and improve the application of federated learning. Our work can lead to massive reduction in communication overheads while still being robust to class imbalance issue. To address these issues more, in future we plan to work on client sampling and investigate the use of heterogeneous models in federated learning.

References

- [1] Ahn, J.H., Kim, G.Y., Kim, D.H., You, C.: Model compression by count-sketch for over-the-air stateless federated learning. IEEE Internet of Things Journal pp. 1–1 (2024).

- 265 <https://doi.org/10.1109/JIOT.2024.3376771>
- 266 [2] Caldas, S., Konečný, J., McMahan, H.B., Talwalkar, A.: Expanding the reach of federated
267 learning by reducing client resource requirements. *arXiv preprint arXiv:1812.07210* (2018)
- 268 [3] Chen, L., Liu, W., Chen, Y., Wang, W.: Communication-efficient design for quantized decen-
269 tralized federated learning. *IEEE Transactions on Signal Processing* (2024)
- 270 [4] Cui, Y., Jia, M., Lin, T.Y., Song, Y., Belongie, S.: Class-balanced loss based on effective number
271 of samples. pp. 9260–9269 (06 2019). <https://doi.org/10.1109/CVPR.2019.00949>
- 272 [5] Elkordy, A.R., Avestimehr, A.S.: Heterosag: Secure aggregation with heterogeneous quantiza-
273 tion in federated learning. *IEEE Transactions on Communications* **70**(4), 2372–2386 (2022)
- 274 [6] Feng, Z., Chen, X., Wu, Q., Wu, W., Zhang, X., Huang, Q.: Feddd: Toward communication-
275 efficient federated learning with differential parameter dropout. *IEEE Transactions on Mobile*
276 *Computing* (2023)
- 277 [7] Gal, Y., Ghahramani, Z.: A theoretically grounded application of dropout in recurrent neural
278 networks. *Advances in neural information processing systems* **29** (2016)
- 279 [8] Gu, X., Sabrina, F., Fan, Z., Sohail, S.: A review of privacy enhancement methods for fed-
280 erated learning in healthcare systems. *International Journal of Environmental Research and*
281 *Public Health* **20**(15) (2023). <https://doi.org/10.3390/ijerph20156539>, <https://www.mdpi.com/1660-4601/20/15/6539>
- 283 [9] Horvath, S., Laskaridis, S., Almeida, M., Leontiadis, I., Venieris, S., Lane, N.: Fjord: Fair and
284 accurate federated learning under heterogeneous targets with ordered dropout. *Advances in*
285 *Neural Information Processing Systems* **34**, 12876–12889 (2021)
- 286 [10] Hospedales, T., Antoniou, A., Micaelli, P., Storkey, A.: Meta-learning in neural networks: A
287 survey. *IEEE transactions on pattern analysis and machine intelligence* **44**(9), 5149–5169 (2021)
- 288 [11] Huisman, M., Van Rijn, J.N., Plaat, A.: A survey of deep meta-learning. *Artificial Intelligence*
289 *Review* **54**(6), 4483–4541 (2021)
- 290 [12] Jacob, B., Kligys, S., Chen, B., Zhu, M., Tang, M., Howard, A., Adam, H., Kalenichenko, D.:
291 Quantization and training of neural networks for efficient integer-arithmetic-only inference. In:
292 *Proceedings of the IEEE conference on computer vision and pattern recognition*. pp. 2704–2713
293 (2018)
- 294 [13] Jeon, I., Hong, M., Yun, J., Kim, G.: Federated learning via meta-variational dropout. *Advances*
295 *in Neural Information Processing Systems* **36** (2024)
- 296 [14] Jia, Z., Jafar, S.A.: X-secure t-private federated submodel learning with elastic dropout resilience.
297 *IEEE Transactions on Information theory* **68**(8), 5418–5439 (2022)
- 298 [15] Konečný, J., McMahan, H.B., Yu, F.X., Richtárik, P., Suresh, A.T., Bacon, D.: Federated
299 learning: Strategies for improving communication efficiency. *arXiv preprint arXiv:1610.05492*
300 (2016)
- 301 [16] Krizhevsky, A., Hinton, G., et al.: Learning multiple layers of features from tiny images (2009)
- 302 [17] Lang, N., Sofer, E., Shaked, T., Shlezinger, N.: Joint privacy enhancement and quantization in
303 federated learning. *IEEE Transactions on Signal Processing* **71**, 295–310 (2023)
- 304 [18] Li, K., Xiao, C.: Pbfl: Communication-efficient federated learning via parameter predicting.
305 *The Computer Journal* **66**(3), 626–642 (2021). <https://doi.org/10.1093/comjnl/bxab184>
- 306 [19] Li, L., Duan, M., Liu, D., Zhang, Y., Ren, A., Chen, X., Tan, Y., Wang, C.: Fedsae: A novel
307 self-adaptive federated learning framework in heterogeneous systems. In: *2021 International*
308 *Joint Conference on Neural Networks (IJCNN)*. pp. 1–10. IEEE (2021)

- [20] Li, T., Sahu, A.K., Talwalkar, A., Smith, V.: Federated learning: Challenges, methods, and future directions. *IEEE Signal Processing Magazine* **37**(3), 50–60 (May 2020). <https://doi.org/10.1109/msp.2020.2975749>, <http://dx.doi.org/10.1109/MSP.2020.2975749>
- [21] Li, Y., Li, W., Xue, Z.: Federated learning with stochastic quantization. *International Journal of Intelligent Systems* **37**(12), 11600–11621 (2022)
- [22] Lim, W.Y.B., Ng, J.S., Xiong, Z., Jin, J., Zhang, Y., Niyato, D., Leung, C., Miao, C.: Decentralized edge intelligence: A dynamic resource allocation framework for hierarchical federated learning. *IEEE Transactions on Parallel and Distributed Systems* **33**(3), 536–550 (2021)
- [23] Lin, W.C., Tsai, C.F., Hu, Y.H., Jhang, J.S.: Clustering-based undersampling in class-imbalanced data. *Information Sciences* **409**, 17–26 (2017)
- [24] Mao, Y., Zhao, Z., Yan, G., Liu, Y., Lan, T., Song, L., Ding, W.: Communication-efficient federated learning with adaptive quantization. *ACM Transactions on Intelligent Systems and Technology (TIST)* **13**(4), 1–26 (2022)
- [25] McMahan, H.B., Moore, E., Ramage, D., y Arcas, B.A.: Federated learning of deep networks using model averaging. *CoRR* **abs/1602.05629** (2016), <http://arxiv.org/abs/1602.05629>
- [26] McMahan, H.B., Ramage, D., Talwar, K., Zhang, L.: Learning differentially private language models without losing accuracy. *CoRR* **abs/1710.06963** (2017), <http://arxiv.org/abs/1710.06963>
- [27] Mills, J., Hu, J., Min, G.: Multi-task federated learning for personalised deep neural networks in edge computing. *IEEE Transactions on Parallel and Distributed Systems* **33**(3), 630–641 (2021)
- [28] Nader, B., Jiahui, H., Hui, Z., Xin, L.: Adaptive federated dropout: Improving communication efficiency and generalization for federated learning. *arXiv preprint arXiv:2011.04050* (2020)
- [29] Niu, Z., Dong, H., Qin, A.: Fedspu: Personalized federated learning for resource-constrained devices with stochastic parameter update. *arXiv preprint arXiv:2403.11464* (2024)
- [30] Ogundokun, R.O., Misra, S., Maskeliunas, R., Damasevicius, R.: A review on federated learning and machine learning approaches: Categorization, application areas, and blockchain technology. *Information* **13**(5) (2022), <https://www.mdpi.com/2078-2489/13/5/263>
- [31] Oh, Y., Jeon, Y.S., Chen, M., Saad, W.: Fedvqcs: Federated learning via vector quantized compressed sensing. *IEEE Transactions on Wireless Communications* (2023)
- [32] Ozkara, K., Singh, N., Data, D., Diggavi, S.: Quped: Quantized personalization via distillation with applications to federated learning. *Advances in Neural Information Processing Systems* **34**, 3622–3634 (2021)
- [33] Sakho, A., Scornet, E., Malherbe, E.: Theoretical and experimental study of smote: limitations and comparisons of rebalancing strategies. *ArXiv* **abs/2402.03819** (2024), <https://api.semanticscholar.org/CorpusID:267499706>
- [34] Shao, J., Sun, Y., Li, S., Zhang, J.: Dres-fl: Dropout-resilient secure federated learning for non-iid clients via secret data sharing. *Advances in Neural Information Processing Systems* **35**, 10533–10545 (2022)
- [35] Shen, Z., Cervino, J., Hassani, H., Ribeiro, A.: An agnostic approach to federated learning with class imbalance. In: *International Conference on Learning Representations* (2021)
- [36] Shlezinger, N., Chen, M., Eldar, Y.C., Poor, H.V., Cui, S.: Uveqfed: Universal vector quantization for federated learning. *IEEE Transactions on Signal Processing* **69**, 500–514 (2020)
- [37] Singh, S.K., Bhardwaj, S.H., Pandey, H., Beniwal, G.: Anomaly detection using federated learning. *Proceedings of International Conference on Artificial Intelligence and Applications* (2020), <https://api.semanticscholar.org/CorpusID:224850507>

- 355 [38] Singla, S., Singla, S., Feizi, S.: Improved deterministic l2 robustness on cifar-10 and cifar-100.
356 arXiv preprint arXiv:2108.04062 (2021)
- 357 [39] Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., Salakhutdinov, R.: Dropout: a simple
358 way to prevent neural networks from overfitting. *The journal of machine learning research* **15**(1),
359 1929–1958 (2014)
- 360 [40] Tang, Y., Liang, Y., Liu, Y., Zhang, J., Ni, L., Qi, L.: Reliable federated learning based on dual-
361 reputation reverse auction mechanism in internet of things. *Future Generation Computer Systems*
362 **156**, 269–284 (2024). <https://doi.org/https://doi.org/10.1016/j.future.2024.03.019>, <https://www.sciencedirect.com/science/article/pii/S0167739X2400089X>
- 364 [41] Team, I.G.P.: Eu general data protection regulation (gdpr)—an implementation and compliance
365 guide. IT Governance Ltd (2020)
- 366 [42] Wang, I., Nair, P., Mahajan, D.: Fluid: Mitigating stragglers in federated learning using invariant
367 dropout. *Advances in Neural Information Processing Systems* **36** (2024)
- 368 [43] Wen, D., Jeon, K.J., Huang, K.: Federated dropout—a simple approach for enabling federated
369 learning on resource constrained devices. *IEEE wireless communications letters* **11**(5), 923–927
370 (2022)
- 371 [44] Xu, J., Du, W., Jin, Y., He, W., Cheng, R.: Ternary compression for communication-efficient
372 federated learning. *IEEE Transactions on Neural Networks and Learning Systems* **33**(3), 1162–
373 1176 (2020)
- 374 [45] Xue, J., Liu, M., Sun, S., Wang, Y., Jiang, H., Jiang, X.: Fedbiad: Communication-efficient
375 and accuracy-guaranteed federated learning with bayesian inference-based adaptive dropout. In:
376 2023 IEEE International Parallel and Distributed Processing Symposium (IPDPS). pp. 489–500.
377 IEEE (2023)
- 378 [46] Zeng, Y., Mu, Y., Yuan, J., Teng, S., Zhang, J., Wan, J., Ren, Y., Zhang, Y.: Adaptive Fed-
379 erated Learning With Non-IID Data. *The Computer Journal* **66**(11), 2758–2772 (09 2022).
380 <https://doi.org/10.1093/comjnl/bxac118>, <https://doi.org/10.1093/comjnl/bxac118>
- 381 [47] Zhang, L., Xu, J., Vijayakumar, P., Sharma, P.K., Ghosh, U.: Homomorphic encryption-based
382 privacy-preserving federated learning in iot-enabled healthcare system. *IEEE Transactions on*
383 *Network Science and Engineering* (2022)